



Ron Davis

AWS MARKETPLACE GTM

SESSION 4 · FUTURE-PROOF

Get bought when the buyer is an agent.

AI agents are starting to shop, evaluate, and transact. Today we make your product discoverable and purchasable when no human is in the loop.

1:1 GTM Coaching · Ron Davis · From listed to bought.

A new buyer just **showed up.**

Procurement research, vendor comparison, and increasingly the purchase itself are moving to AI agents. They don't read your homepage — they parse your data and match it to a requirement.

// IT SHOPS

Discovers

An agent gathers candidate vendors from listings, docs, and the catalog — by requirement, not brand.

// IT JUDGES

Evaluates

It compares specs, units, and terms against the buyer's constraints — and silently drops what it can't parse.

// IT ACTS

Transacts

Where the path is machine-readable, it can move to purchase without waiting on a human handshake.

// THE STAKES

An agent doesn't ask follow-up questions. Ambiguity isn't a conversation — it's a disqualification.

Every gap a human would forgive, an agent treats as a "no." Agent-ready is just clarity, made literal.

A machine-legible **product statement**.

Strip the marketing voice. State what it is, the job it does, and its inputs and outputs — in plain nouns an agent can match to a requirement.

```
product : what it is, in one literal noun phrase
job      : the task it completes for the buyer
inputs   : what it takes in (data, scope, prerequisites)
outputs  : what it produces, in measurable terms
limits   : what it does NOT do – so the agent doesn't mis-match
```

"When an agent is doing X, you're the right tool."

An agent picks tools by the task in front of it. Trigger tasks are the specific jobs where you should be the obvious choice — named so explicitly an agent reaches for you.

- 1 **Name the task, not your category** — "when reconciling multi-cloud spend," not "FinOps platform."
- 2 **Make it matchable** — phrase it the way the requirement would arrive: a verb plus an object.
- 3 **Cover the top three jobs** you uniquely win — depth beats a long, vague list.

// WHY IT MATTERS

Agents select on the job-to-be-done. If you're not described as a job, you're not in the running.

The trigger-task list is the bridge between what you do and what an agent is actually trying to accomplish.

Write your trigger-task list.

We draft the three tasks where an agent should choose you — phrased as the requirement would arrive, not as your pitch.

#	"WHEN AN AGENT IS DOING..."	"...REACH FOR US BECAUSE"
01	-----	-----
02	-----	-----
03	-----	-----

Is everything structured and parseable?

We walk your listing, docs, and pricing and ask one question of each surface: can a machine extract clean facts, or is it prose it has to guess at?

// SURFACE ONE

Listing

Capabilities, limits, and use-cases in consistent structure — not one paragraph an agent must summarize.

// SURFACE TWO

Docs

Inputs, outputs, integrations, and constraints stated literally and findable — the agent's evidence layer.

// SURFACE THREE

Pricing

Expressible as data: units, rates, and a base path the agent can compute without a call.

Where an agent **trips**.

Humans fill gaps with assumptions; agents stall on them. We hunt down the specific places your content forces a guess — and replace the guess with a fact.

- ! **Vague capability claims** — "powerful," "flexible," "enterprise-grade" carry zero matchable meaning.
- ! **Implied units** — pricing or limits that a human infers but an agent can't extract.
- ! **Missing limits** — no statement of what you don't do, so the agent mis-matches you to the wrong job.
- ! **Conflicting facts** across listing, docs, and pricing — the agent can't tell which to trust.

Terms an agent can **accept**.

Being chosen isn't enough if the agent can't complete the purchase. The base path has to be clear units, predictable cost, and terms that don't require a human to sign.

// UNITS

Clear units

A named, defined unit the agent reads and reasons about — not a bespoke quote.

// COST

Predictable cost

Computable from inputs before commit, so the agent can budget and decide.

// TERMS

Acceptable terms

Standard contract or metered path the agent can accept without a handshake.

// THE BRIDGE TO SESSION 3

Agent-ready pricing and a procurement-defensible price are the same work.

The self-serve, legible path you built last session is exactly what makes you transactable here.

Your agent-ready checklist, **by impact.**

- | | | |
|---|--|---------|
| 1 | Machine-legible product statement — what / job / inputs / outputs / limits, published. | HIGHEST |
| 2 | Structured listing description — capabilities and use-cases in consistent, scannable structure. | HIGHEST |
| 3 | Deterministic base pricing — a self-serve path expressible as data, no "contact us." | HIGH |
| 4 | Trigger-task list — the three jobs where an agent should pick you, phrased to match. | HIGH |
| 5 | Resolve ambiguity & conflicts — align facts across listing, docs, and pricing. | MEDIUM |
| 6 | Machine-consumable specs in docs — inputs, outputs, integrations, constraints, literal. | MEDIUM |

Apply the top three.

The checklist is ordered for a reason — the first three fixes carry most of the gain. Ship those this week; the rest follow.

- ☐ Publish the machine-legible product statement. — what / job / inputs / outputs / limits.
- ☐ Restructure the listing description into scannable structure. — no more prose blob.
- ☐ Make the base pricing path deterministic and self-serve. — an agent can transact.
- ☐ Re-run Agent-Ready Positioning to test the new copy. — verify it parses cleanly.